

# Recap

## Chapter 1: Parallel computing introduction

Zhiming Zhao  
University of Amsterdam

	Shared memory			Distributed memory		
	Single core	Multi core	Many cores...	Cluster	Cloud	Blockchain...
Sequential application						
Concurrent application						
Parallel application						
Distributed application						

# What is the difference between concurrency and parallelism?

- Quiz: L1Q1
- DPP2025

	Shared memory			Distributed memory		
	Single core	Multi core	Many cores...	Cluster	Cloud	Blockchain...
Sequential application	<i>Can be executed on all machines, but may not fully utilizing the resources.</i>					
Concurrent application	<i>Can be executed</i>					
Parallel application		<i>Can be executed</i>				
Distributed application				<i>Can be executed</i>		

# Chapter 2

# Introduction to multithreading and

# PThread

Zhiming Zhao

Multiscale Networked Systems

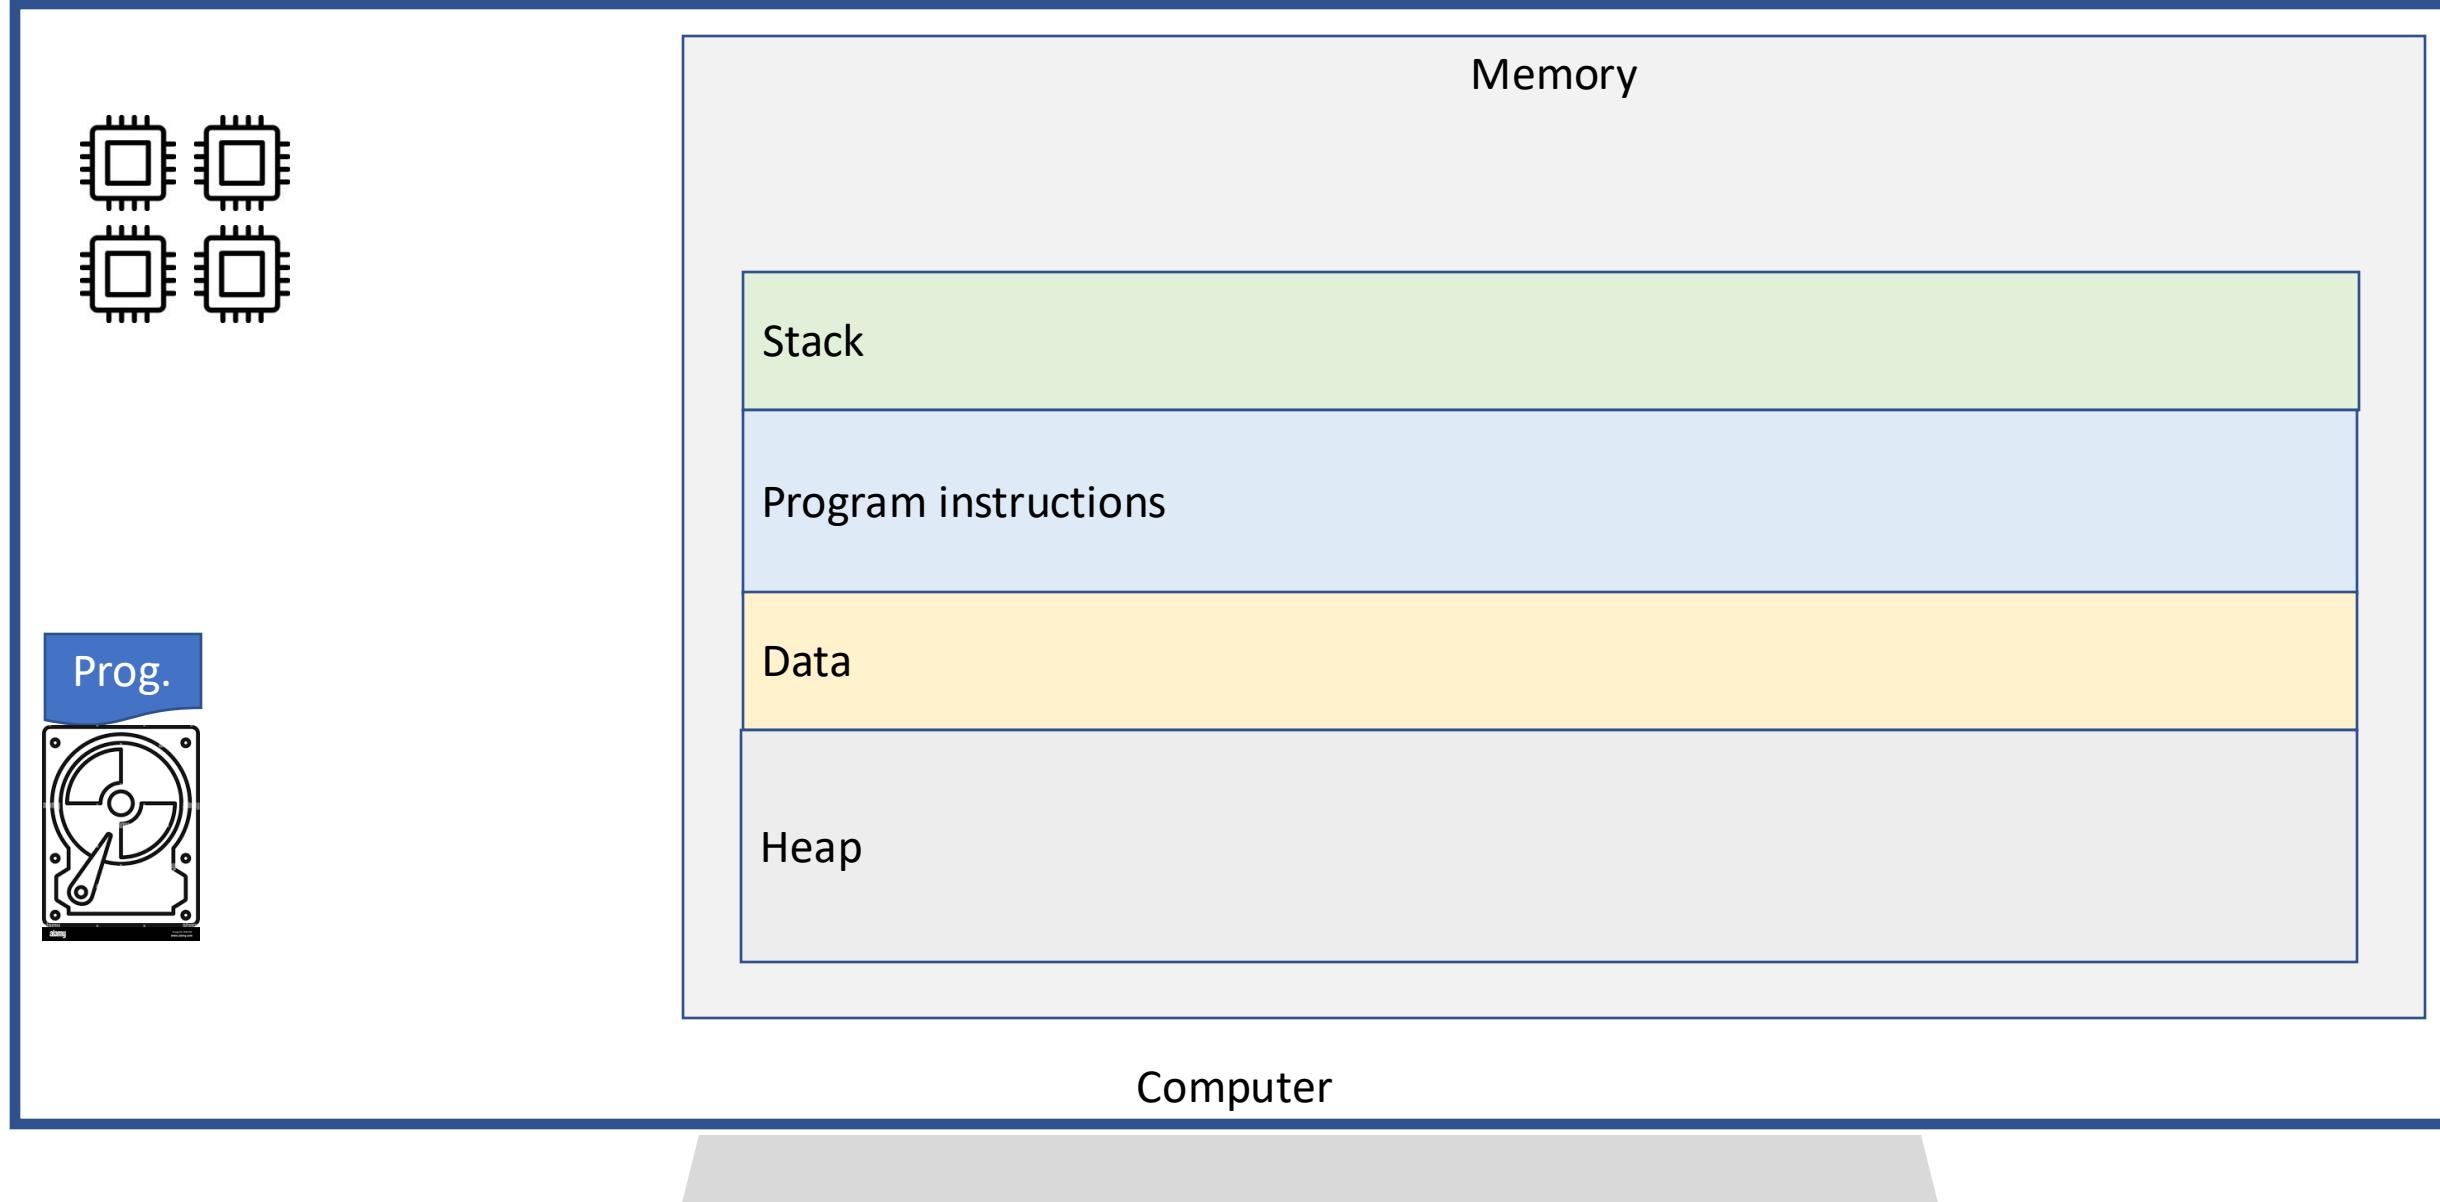
University of Amsterdam

# Content

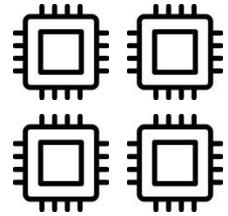
1. Multithreading
2. pThread
3. Race condition
4. Data parallelisation using pThread
5. Speedup and efficiency
6. Task parallelization

# 1. Multithreading

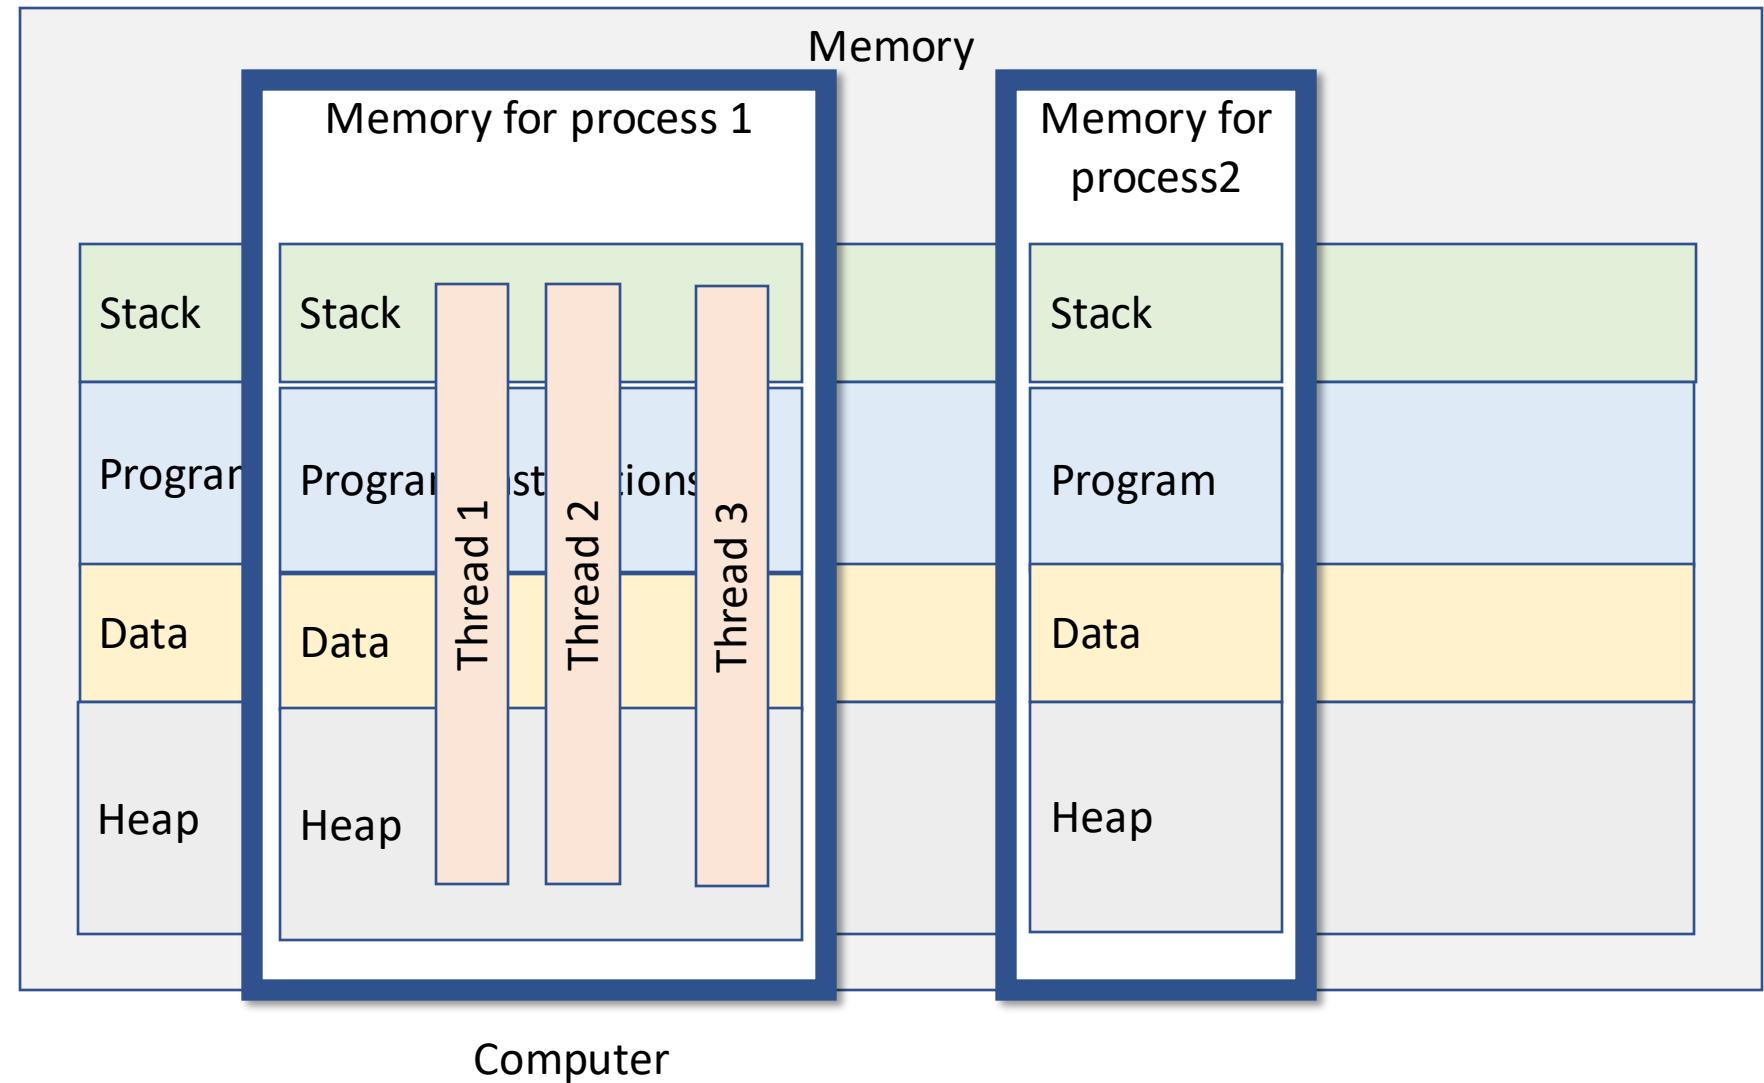
# Process and thread



# Process and thread



Being executed by the **Operating System**, a **program** will be loaded in the memory as a **process**.



Threads are “mini processes” inside a process, and share the memory of the process.

# Process and Threads

- Processes
  - control their own resources
  - No shared memory
  - Communication through messages or files
  - Heavy-weight (expensive to set up and destroy)
- Threads
  - Control their limited local memory and shared address space
  - Communication through shared memory
  - Light-weight (much cheaper to set up and destroy)

# Programming models

- Pthreads + intrinsics
- TBB – Thread building blocks
  - Threading library
- OpenCL, CUDA
  - To be discussed ...
- OpenMP
  - High-level, pragma-based
- ... <many more>
- Cilk
  - Simple divide-and-conquer model
- ... <many more>

L03

L04



Level of abstraction increases

# What will we learn?

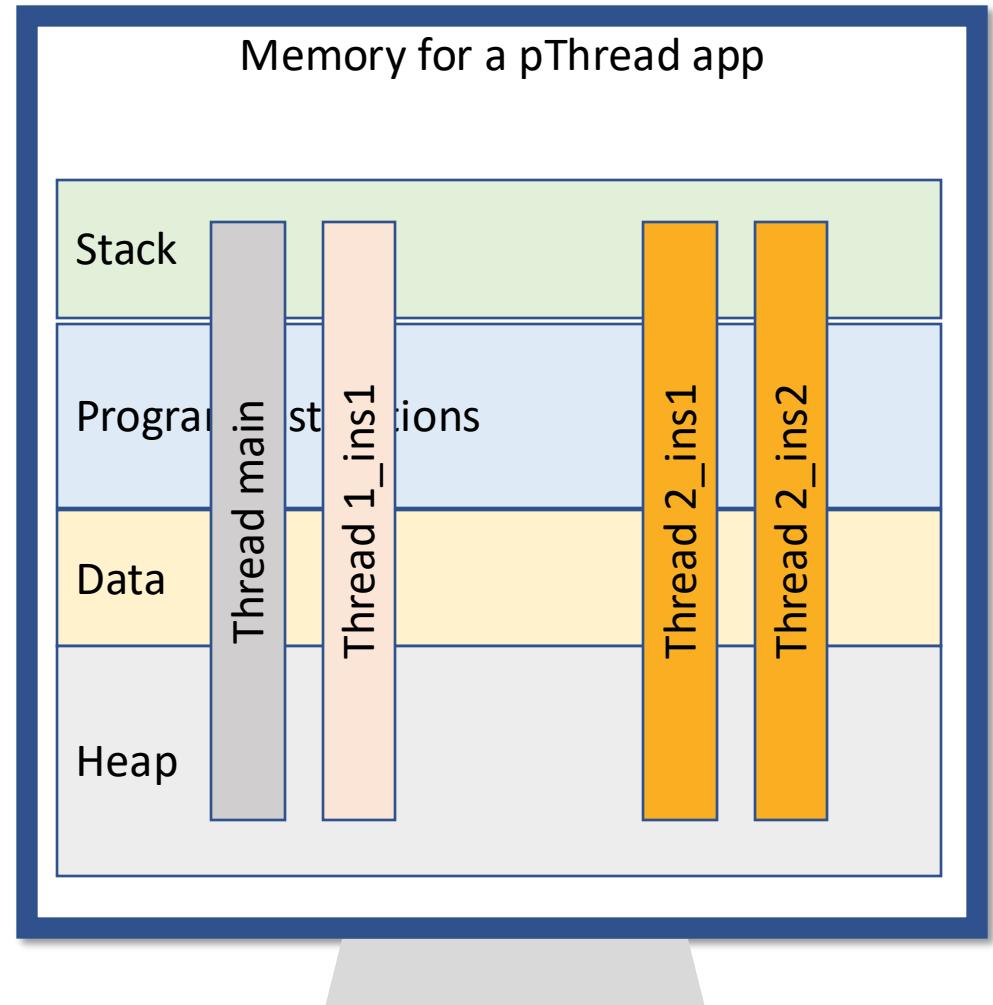
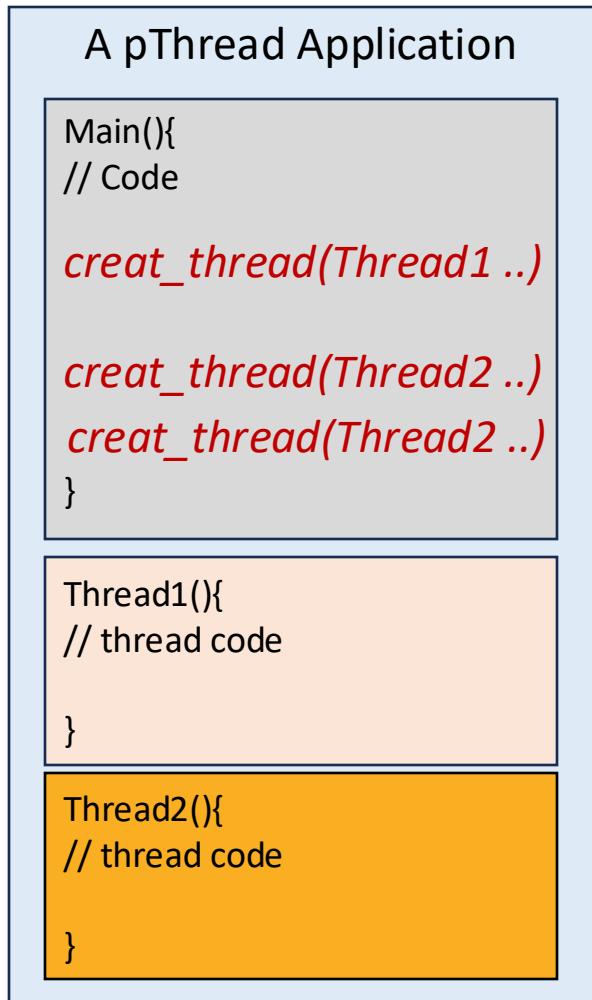
- How to create a thread?
- How to perform computations within a thread?
- How do threads communicate with each other?
- How do threads concurrently access shared memory?
- How can you stop or terminate a thread?
- How do you design an appropriate multithreaded program for a specific problem?
- How do we evaluate the performance of threads?

## 2. Multithreading with POSIX threads

# What are “Pthreads”

- Standardized programming interface for threads
  - UNIX systems : IEEE POSIX (Portable Operating Systems Interface) 1003.1c standard (1995)
- All implementations of these standards are known as **POSIX threads** or **threads**.
- Pthreads are defined as a set of **C language programming types** and **procedure called**
  - pthread.h header/include file
  - a thread library
    - Or part of libc in some implementations.

# A high level view of multi thread program



# Create and start a thread

```
int pthread_create (  
    pthread_t* thread, // OUT: thread object  
    pthread_attr_t* attributes, // IN: attributes  
    void* (*start_routine)(void *), // IN: routine to start thread  
    void* argument// IN: input attributes to the start routine  
);
```

## 2.1 Create and start a thread

```
int pthread_create (
    pthread_t* thread, // OUT: thread object
    pthread_attr_t* attributes, // IN: attributes
    void* (*start_routine)(void *), // IN: routine to start thread
    void* argument); // IN: input attributes to the start routine
```

```
int pthread_join (
    pthread_t thread_id , // IN: thread object
    void **return_value ) // OUT: return value
```

# A “hello world” example

```
/*
* hello.c - Pthreads "hello, world" program
*/
#include <pthread.h>
void *thread(void *vargp);
int main()
{
    pthread_t thd;
    pthread_create(&thd, NULL, thread, NULL);
    pthread_join(thd, NULL);
    exit(0);
}

void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

1. #include <pthread.h>
2. pthread\_create(&thd, NULL, thread, NULL);
3. pthread\_join(thd, NULL);

# To create a thread

```
/*
 * hello.c - Pthreads "hello, world" program
 */

#include <pthread.h>

void *thread(void *vargp);

int main()
{
    pthread_t thd;
    pthread_create(&thd, NULL, thread, NULL);
    pthread_join(thd, NULL);
    exit(0);
}

void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

1. Thread routine name
2. Thread object

Question: 1) how many threads will be created? 2) what will be the output?

# To create multiple threads

```
/*
 * hello.c - Pthreads "hello, world" program
 */

#include <pthread.h>

void *thread(void *vargp);

int main()
{ int number=5;
  pthread_t *thd;
  thd = (pthread_t*) malloc (sizeof(pthread_t)*number);

  for (int i=0;i<number; i++)
    pthread_create( thd[i], NULL, thread, NULL);
  for (int i=0;i<number;i++)
    pthread_join(thd[i], NULL);
  exit(0);
}

void *thread(void *vargp) /* thread routine */
{
  printf("Hello,world!\n");
  return NULL;
}
```

1. Each thread needs a dedicated **pthread\_t object**
2. Create thread **one by one**, they can have **different thread routine**
3. Put the **join** after creating all threads.

## Discussion:

1. Do you have to create threads in the **main** function?
2. Any bug you can find in the code?

# To create multiple threads

```
/*
 * hello.c - Pthreads "hello, world" program
 */

#include <pthread.h>

void *thread(void *vargp);

int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);

    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    exit(0);
}

void *thread(void *vargp) /* thread routine */
{
    printf("Hello,world!\n");
    return NULL;
}
```

1. Each thread needs a dedicated **pthread\_t object**
2. Create thread **one by one**, they can have **different thread routine**
3. Put the **join** after creating all threads.
4. **Pthread\_t \*** in the create function!!!

Question: How does a thread know its position among all threads?

# Another example

```
#include <pthread.h>

void *thread(void *vargp);

int count=0;

int main()
{
    int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);

    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}

void *thread(void *vargp) /* thread routine */
{
    for (int i=0;i< 100;i++)
        count++;
    return NULL;
}
```

## Discussion

1. What will the program print?
  1. 500
  2.  $\geq 500$
  3.  $> 500$
  4.  $\leq 500$
  5.  $< 500$
  6. 0

2. How to solve it?

count ++ will be executed as i) read count, ii) add count with 1, iii) write count

# Variables and memory access

```
#include <pthread.h>

void *thread(void *vargp);

int count=0;

int main()
{
    int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);

    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}

void *thread(void *vargp) /* thread routine */
{
    static int total = 100;
    for (int i=0;i< total;i++)
        count++;
    return NULL;
}
```

- Global variables

- Def: Variable declared outside of any function
- 

- Local variables

- Def: Variables declared inside functions, without the **static** attribute
- 

- Local static variables

- Def: Variable declared inside functions, with the **static** attribute
- 

Question: What are the global, local, and local static variables?

# Variables and memory access

```
#include <pthread.h>

void *thread(void *vargp);

int count=0;
int main()
{ int number=5;
  pthread_t *thd;
  thd = (pthread_t*) malloc (sizeof(pthread_t)*number);

    for (int i=0;i<number; i++)
      pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
      pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}

void *thread(void *vargp) /* thread routine */
{
  static int total = 100;
  for (int i=0;i< total;i++)
    count++;
  return NULL;
}
```

- Global variables

- Def: Variable declared outside of any function
- Exactly **one instance** exists for each global variable (shared across all functions and threads)

- Local variables

- Def: Variables declared inside functions, without the **static** attribute
- **Each thread stack contains its own instance of each local variable**

- Local static variables

- Def: Variable declared inside functions, with the **static** attribute
- Exactly **one instance** exists for each local static variable, shared across all function calls and threads.

Question: What are the global, local, and local static variables?

## 2.2 Race Condition

- A race condition occurs when two or more threads in one application:
  - try to access the same memory location(**variables concurrently**, and
  - more than one of the accesses are for **writing** and
  - the threads **do not use** exclusive mechanisms to **control their access** to that memory.

### 3. Race condition

# Race Condition

A race condition occurs when two or more threads in a single process:

- try to access and change the same variable at the same time
- the threads do not use exclusive mechanisms to control their access

Most solutions are based on

1. Replication
2. Atomic operation
3. Locking

# Choice 1: move global variable to local

```
#include <pthread.h>
void *thread(void *vargp);
int countT[5];
int count=0;

int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, &countT[i]);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    for (int i=0;i<number;i++)
        count+=countT[i];
    printf("Result:%d\n", count);
    exit(0);
}

void *thread(void *vargp) /* thread routine */
{
    int *localCount=(int*)vargp;
    *localCount=0;
    for (int i=0;i< 100;i++)
        (*localCount)++;
    return NULL;
}
```

## Replicate variables

1. Pass variables to each thread: **countT[i]**
2. Each thread adds its own copy of the variable: **localCount**
3. The main thread calculates the final result: **count**

Question: What output will you get?

- 500

## Note:

- How do you pass value or address to the thread?
- Don't forget to initialise the variable.

# Choice 2: Using atomic operation

```
#include <pthread.h>
#include <stdatomic.h>
void *thread(void *vargp);
atomic_int account=ATOMIC_VAR_INIT(0);
int count=0;
int main()
{ int number=5;
  pthread_t *thd;
  thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
  for (int i=0;i<number; i++)
    pthread_create(&thd[i], NULL, thread, NULL);

  for (int i=0;i<number;i++)
    pthread_join(thd[i], NULL);
  printf("Result:%d\n", account);
  exit(0);
}
void *thread(void *vargp) /* thread routine */
{
  for (int i=0;i< 100;i++)
    atomic_fetch_add_explicit( &account, 1, memory_order_relaxed);
  return NULL;
}
```

Solution:

1. Define shared variable as atomic variable: **atomic\_int account**
2. Use atomic operation:  
**atomic\_fetch\_add\_explicit**

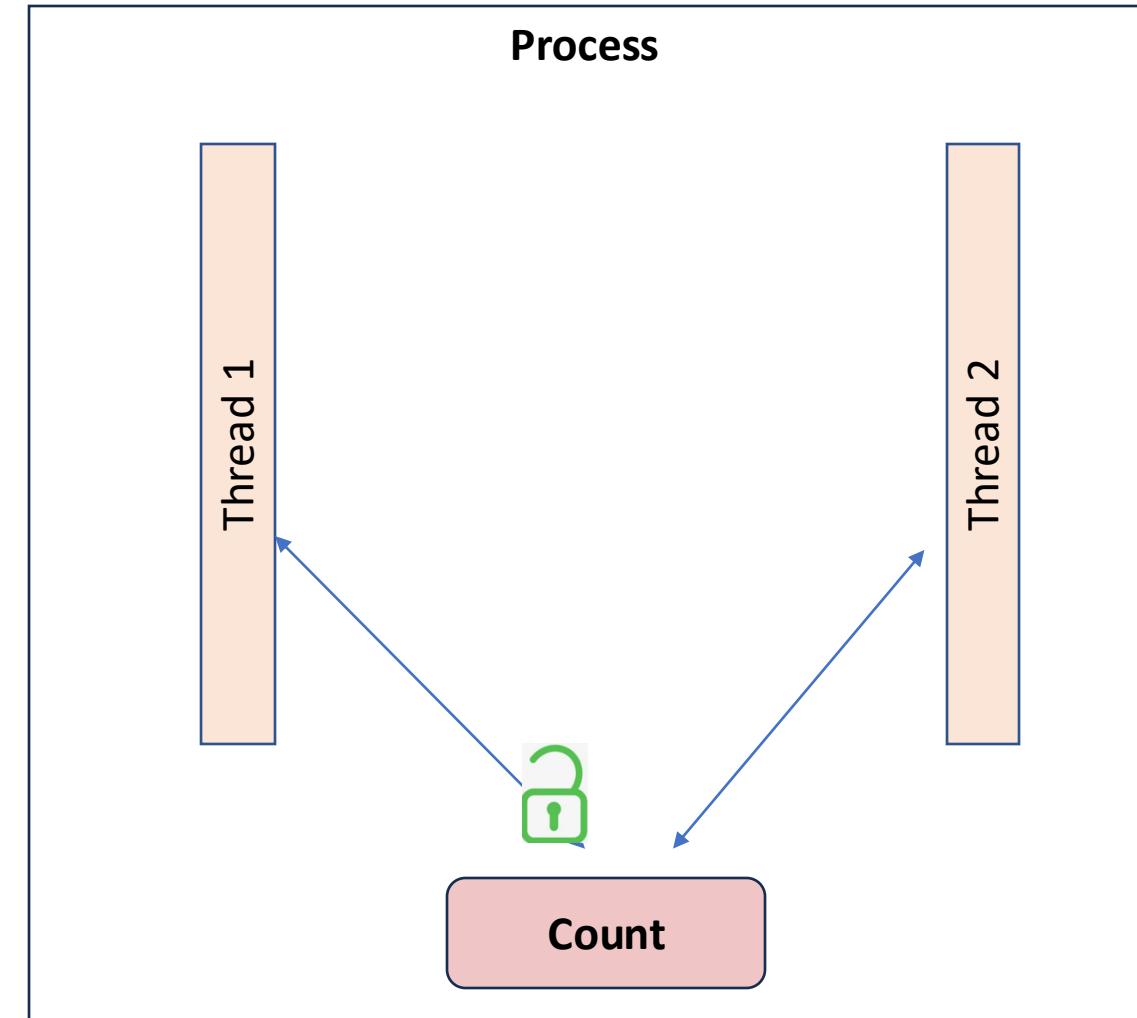
It will also guarantee an output as 500.

Note:

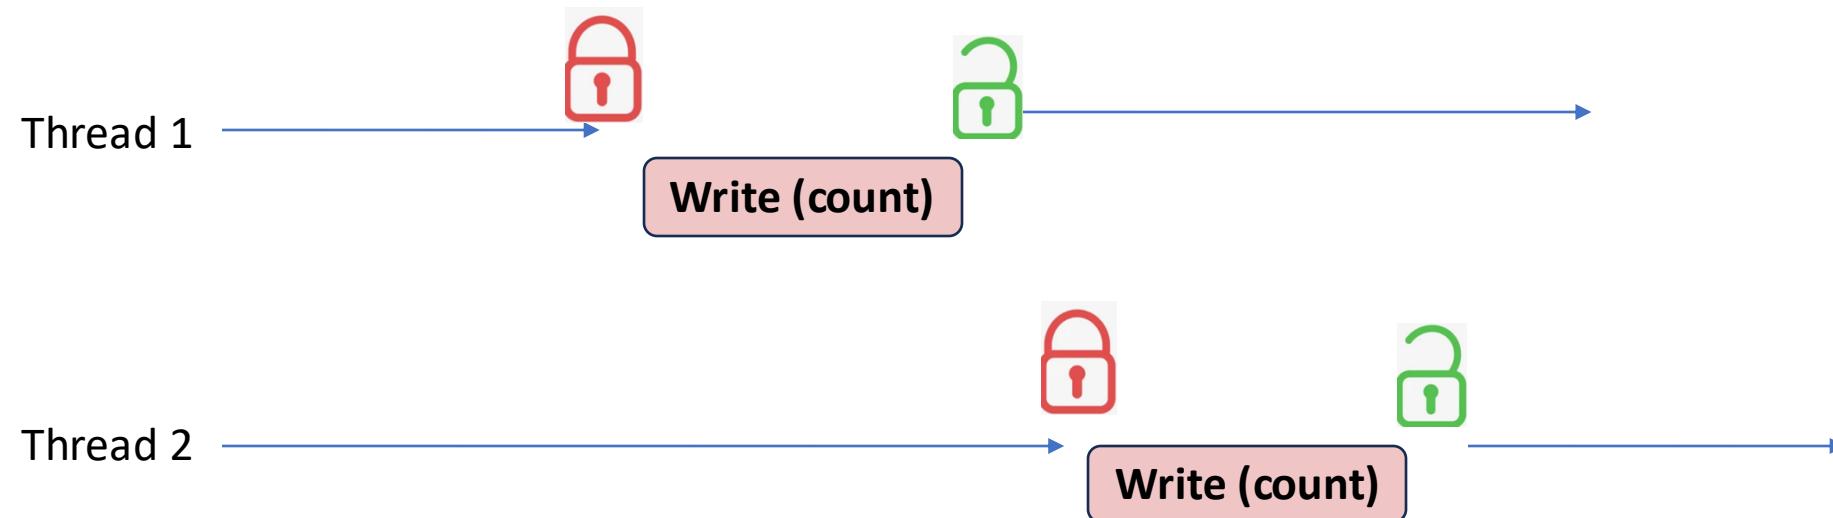
1. Not always possible for complex operations

# Choice 3: exclusive access to critical region

- Threads are NOT allowed to access the same region simultaneously
- Only one thread is allowed to use the region at one time



# How does it work?



1. A thread has to **lock** a critical region if it wants to work on it
2. A thread can only **lock** a critical region when its status is **unlocked** (available)
3. A thread has to **wait** for a locked region to become **unlocked** if it wants to work on it
4. A thread can only enter the critical region when it has **locked** it.
5. A thread has to **unlock** a critical region after it finishes the work on it

# Mutex locks: essential functions

Global/main/static:

```
pthread_mutex_t lock;
```

Main:

```
pthread_mutex_init(&lock, attributes);
```

...

```
pthread_mutex_destroy(&lock);
```

Per-thread:

```
pthread_mutex_lock ( &lock );
// critical region here
pthread_mutex_unlock ( &lock );
```

# Code example

```
#include <pthread.h>
void *thread(void *vargp);
int count=0;
pthread_mutex_t mutex;
int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    pthread_mutex_init(& mutex, NULL);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    pthread_mutex_destroy(& mutex);
    exit(0);
}
void *thread(void *vargp) /* thread routine */
{
    for (int i=0;i< 100;i++)
    { pthread_mutex_lock( &mutex );
        count++;
        pthread_mutex_unlock( &mutex);
    }
    return NULL;
}
```

## Use a lock

1. Critical region: **count**
2. **pthread\_mutex\_t**
3. **pthread\_mutex\_init**
4. **pthread\_mutex\_destroy**
5. **pthread\_mutex\_lock**
6. **pthread\_mutex\_unlock**
7. **pthread\_mutex\_trylock**

Question: What output will you get?

- 500

Note:

- You can define multi mutex objects for lock
- Carefully check which one is using

# More mutex functions

- Static initialization with default attributes:
  - `pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;`
- Dynamic initialization with attributes:
  - `int pthread_mutex_init (pthread_mutex_t *lock,  
                              pthread_mutexattr_t * attributes );`
    - Attributes\* include ERRORCHECK, RECURSIVE, DEFAULT
- Dynamic de-allocation:
  - `int pthread_mutex_destroy ( pthread_mutex_t * lock );`
- Optimistic locking:
  - `int pthread_mutex_trylock ( pthread_mutex_t * lock );`
    - // returns 0 upon success
    - // returns EBUSY upon failure
- (counting) Semaphores: multiple resources => multiple threads
  - `sem_init(), sem_wait(), sem_post(), ...`

\*Check, for example: `pthread_mutexattr_gettype`

# Choice 3: Enforce the critical region

```
#include <pthread.h>
void *thread(void *vargp);
int count=0;
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    pthread_mutex_init(& mutex, NULL);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    pthread_mutex_destroy(& mutex);
    exit(0);
}
void *thread(void *vargp) /* thread routine */
{
    for (int i=0;i< 100;i++)
    { pthread_mutex_lock( &mutex );
        count++;
        pthread_mutex_unlock( &mutex);
    }
    return NULL;
}
```

Where should we put the lock?

```
void *thread(void *vargp) /* thread routine */
{
    pthread_mutex_lock( &mutex );
    for (int i=0;i< 100;i++)
        count++;
    pthread_mutex_unlock( &mutex);
    return NULL;
}
```

# Mutex locks

- Mutex locks are abstract data objects.
- Only one thread at a time may hold a mutex lock.
- Threads block upon locking if lock is unavailable.
- Locking / unlocking are guaranteed to be **atomic**.
- No **fairness** on waiting threads upon unlocking.
- Only the owner of the lock can unlock it.
- **Re-locking by the owner causes deadlock.**

# Question: what is the output?

```
#include <pthread.h>
void *thread(void *vargp);
int count=0;
int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}
void *thread(void *vargp) /* thread routine */
{
    pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;

    for (int i=0;i< 100;i++)
    { pthread_mutex_lock( &mutex );
        count++;
        pthread_mutex_unlock( &mutex );
    }
    return NULL;
}
```

What is the output?

1. 500
2. 0
3. <=500

# Question: what is the output?

```
#include <pthread.h>
void *thread(void *vargp);
int count=0;
int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}
void *thread(void *vargp) /* thread routine */
{
static pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
    for (int i=0;i< 100;i++)
    { pthread_mutex_lock( &mutex );
        count++;
        pthread_mutex_unlock( &mutex);
    }
    return NULL;
}
```

What is the output?

1. 500
2. 0
3. <=500

Be very careful when using static variables in the threads.

# Question: what is the output?

```
#include <pthread.h>
void *thread(void *vargp);
int count=0;
int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}
void *thread(void *vargp) /* thread routine */
{
static pthread_mutex_t mutex;
pthread_mutex_init(&mutex, NULL);
for (int i=0;i< 100;i++)
{ pthread_mutex_lock( &mutex );
    count++;
    pthread_mutex_unlock( &mutex );
}
return NULL;
}
```

What is the output?

1. 500
2. 0
3. <=500

Be very careful when you call a function inside a thread. Make sure it is **thread safe**.

# Locks & Deadlocks

- **Locking** = a thread locks a resource for single use.
- **Deadlock** = threads are blocked, waiting for each other and no progress is being made.
  - Examples?
    - 4-way stop
    - Dining philosophers problem
- **Livelock** = threads “oscillate” between states with regard to one another and no progress is being made.
  - Examples?
    - People meeting “on the corridor”
    - MAC protocol collisions

# Deadlock

Deadlock occurs when

1. Hold and Wait



## Detecting Deadlock:

- Cycle detection, or debugging tools.

## Fixing Deadlock:

- Timeouts with retries, avoid nested locks, use non-blocking lock attempts, or use higher-level concurrency utilities.

Question: How to fix it?

# Summary: race condition solutions

- Solution 1
  - 1 counter per thread
  - Advantage?
    - No data races
  - Disadvantage?
    - More memory
    - Reduction phase
- Solution 2
  - Busy waiting for an atomic operation
  - Advantage?
    - Do not suspend the thread execution.
  - Disadvantage?
    - Not suitable for complex coding blocks
- Solution 3
  - Enforce a critical region
    - Guarantees all operations will execute without interleaving
  - Advantage?
    - Same memory
    - No reduction
  - Disadvantage?
    - Serialization!

Which one is better?

# 4. Data parallelization using pThread

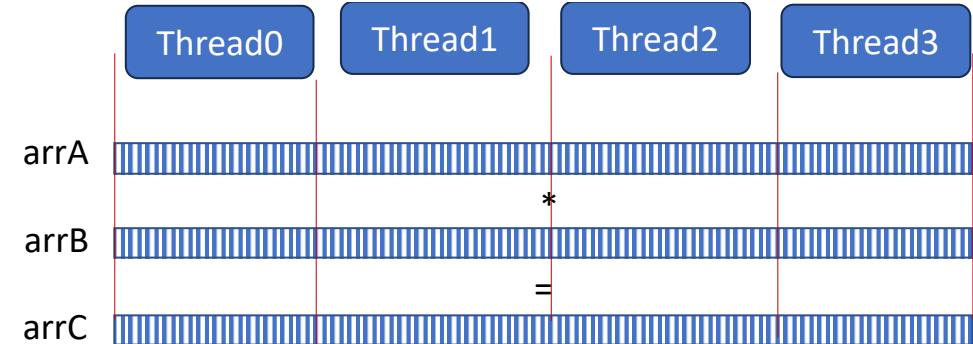
# Discussion: How to do data parallelism?

```
#include <pthread.h>
#include <stdatomic.h>
void *thread(void *vargp);
int *arrA, *arrB, *arrC;
int main()
{
    int number=4;
    int length=100000;
    pthread_t *thd;
    arrA=(int*)malloc(length*sizeof(int));
    arrB=(int*)malloc(length*sizeof(int));
    arrC=(int*)malloc(length*sizeof(int));

    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);

    for (int i=0;i<number; i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", account);
    exit(0);
}

void *thread(void *vargp)
{
    for (int i=?;i< ?;i++)
        // compute arrC[x]=arrA[x]*arrB[x]
    Return NULL;
}
```



Question: How to inform the threads about their tasks?

# How to inform the thread about their task?

```
#include <pthread.h>
#include <stdatomic.h>
void *thread(void *vargp);
int *arrA, *arrB, *arrC;
typedef struct { //data to send to each thread
    int start,end;
    int id;
} t_data;

//create the threads in the main ()
for (i=0; i<number; i++) {
    myData[i].id=i;
    myData[i].start=i*(length/number);
    myData[i].end=(i+1)*(length/number)-1;
    pthread_create(&thd[i], NULL, thread,(void*)&myData[i]);
}
```

```
int main()
{
    int number=5;
    pthread_t *thd;
    t_data *myData;
    // skip the code for initializing thd, myData etc.
    // compute the start/end for each thread
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, (void*)&myData[i]);
    ...
}
```

```
void *thread(void *vargp) /* thread routine */
{
    int start= *((t_data*) vargp).start;
    int end= *((t_data*) vargp).end;
    ...
    for (int i=start ;i< end;i++)
        arrC[x]=arrA[x]*arrB[x]
}
```

- Send it to each thread

# How to synchronize the progress of different threads?

```
int cnt[10000];
int average[NUM_ITER];
void* thread_func(void* args) {
    // obtain start, end from args
    for(int i = 0; i<NUM_ITER; i++) {
        for(int j=start;j<end;j++)
            cnt[j]++;
        //compute_the_average_of_cnt
    }
}

int main() {
    int *tid=malloc(sizeof(int));
    // .. Some code for configuring the start/end for each thread

    for (int i = 0; i < num_threads; i++)
        pthread_create(&threads[i], NULL, thread_func, (void*)&myData);
    for (i = 0; i < num_threads; i++) {
        pthread_join(threads[i], NULL);
    }
    return 0;
}
```

- Threads compute different data chunks often need to be **synchronized**, e.g., from one iteration to the next one
- How do they synchronise the progress?

# Other synchronization primitives

- Barriers: all threads wait at barrier before they all can continue.  
`pthread_barrier_***(...)`  
e.g. `pthread_barrier_init(&barr, attr, count)` `pthread_barrier_wait(&barr)`
- Join: wait for “child” threads to finish and collects results  
`pthread_join(...)`

# How to synchronize the progress of different threads?

```
int cnt[10000];
int NUM_ITER=100;
int num_threads;
pthread_barrier_t barrier;

void* thread_func(void* args) {
    // obtain start, end from args
    for(int i = 0; i<NUM_ITER; i++)
    {
        for(int j=start;j<end;j++)
            cnt[j]++;
        pthread_barrier_wait(&barrier);
        //compute_the_average_of_cnt
    }
}

int main() {
    int *tid=malloc(sizeof(int));
    // .. Some code for configuring the start/end for each thread

    pthread_barrier_init(&barrier, NULL, num_threads);
    for (int i = 0; i < num_threads; i++)
        pthread_create(&threads[i], NULL, thread_func, (void*)&myData);
    for (i = 0; i < num_threads; i++) {
        pthread_join(threads[i], NULL);
    }
    return 0;
}
```

- `pthread_barrier`
- `pthread_barrier_wait`
- `pthread_join`

# Question: what is the output?

```
#include <pthread.h>
void *thread(void *vargp);
int count=0;
int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}
void *thread(void *vargp) /* thread routine */
{
static pthread_mutex_t mutex;
pthread_mutex_init(&mutex, NULL);
for (int i=0;i< 100;i++)
{ pthread_mutex_lock( &mutex );
    count++;
    pthread_mutex_unlock( &mutex );
}
return NULL;
}
```

What is the output?

1. 500
2. 0
3. <=500

Be very careful when you call a function inside a thread. Make sure it is **thread safe**.

# Question: what is the output?

```
#include <pthread.h>
void *thread(void *vargp);
int count=0;
pthread_barrier_t barrier;
int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    pthread_barrier_init(&barrier, NULL, num_threads);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number;i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}
void *thread(void *vargp) /* thread routine */
{
    Static pthread_mutex_t mutex;
    pthread_mutex_init(& mutex, NULL);
    pthread_barrier_wait(&barrier);
    for (int i=0;i< 100;i++)
    { pthread_mutex_lock( &mutex );
        count++;
        pthread_mutex_unlock( &mutex);
    }
    return NULL;
}
```

What is the output?

1. 500
2. 0
3. <=500

# Question: what will be the output?

```
#include <pthread.h>
void *thread(void *vargp);
int count=0;
pthread_barrier_t barrier;
int main()
{ int number=5;
    pthread_t *thd;
    thd = (pthread_t*) malloc (sizeof(pthread_t)*number);
    pthread_barrier_init(&barrier, NULL, num_threads);
    for (int i=0;i<number; i++)
        pthread_create(&thd[i], NULL, thread, NULL);
    for (int i=0;i<number; i++)
        pthread_join(thd[i], NULL);
    printf("Result:%d\n", count);
    exit(0);
}
void *thread(void *vargp) /* thread routine */
{
    Static pthread_mutex_t mutex;
    pthread_mutex_init(& mutex, NULL);
    pthread_barrier_wait(&barrier);
    for (int i=0;i< 100;i++)
    { pthread_mutex_lock( &mutex );
        count++;
        pthread_barrier_wait(&barrier);
        pthread_mutex_unlock( &mutex);
    }
    return NULL;
}
```

What is the output?

1. 500
2. 0
3. <=500
4. Deadlock

# Pthread parallelism

- Data parallelism, e.g., Processing large array
  - Allocate data ranges for each thread
  - Identify potential race conditions
  - Apply lock and synchronization mechanisms

# 5. Speedup and efficiency for data parallelisation

# Speed-up given n processors(cores)

- Highly dependent on Time(1)
- *Best performed sequential execution as the baseline*

$$Speedup(n) = \frac{Time_{best}(1)}{Time(n)}$$

Sequential execution

Parallel execution

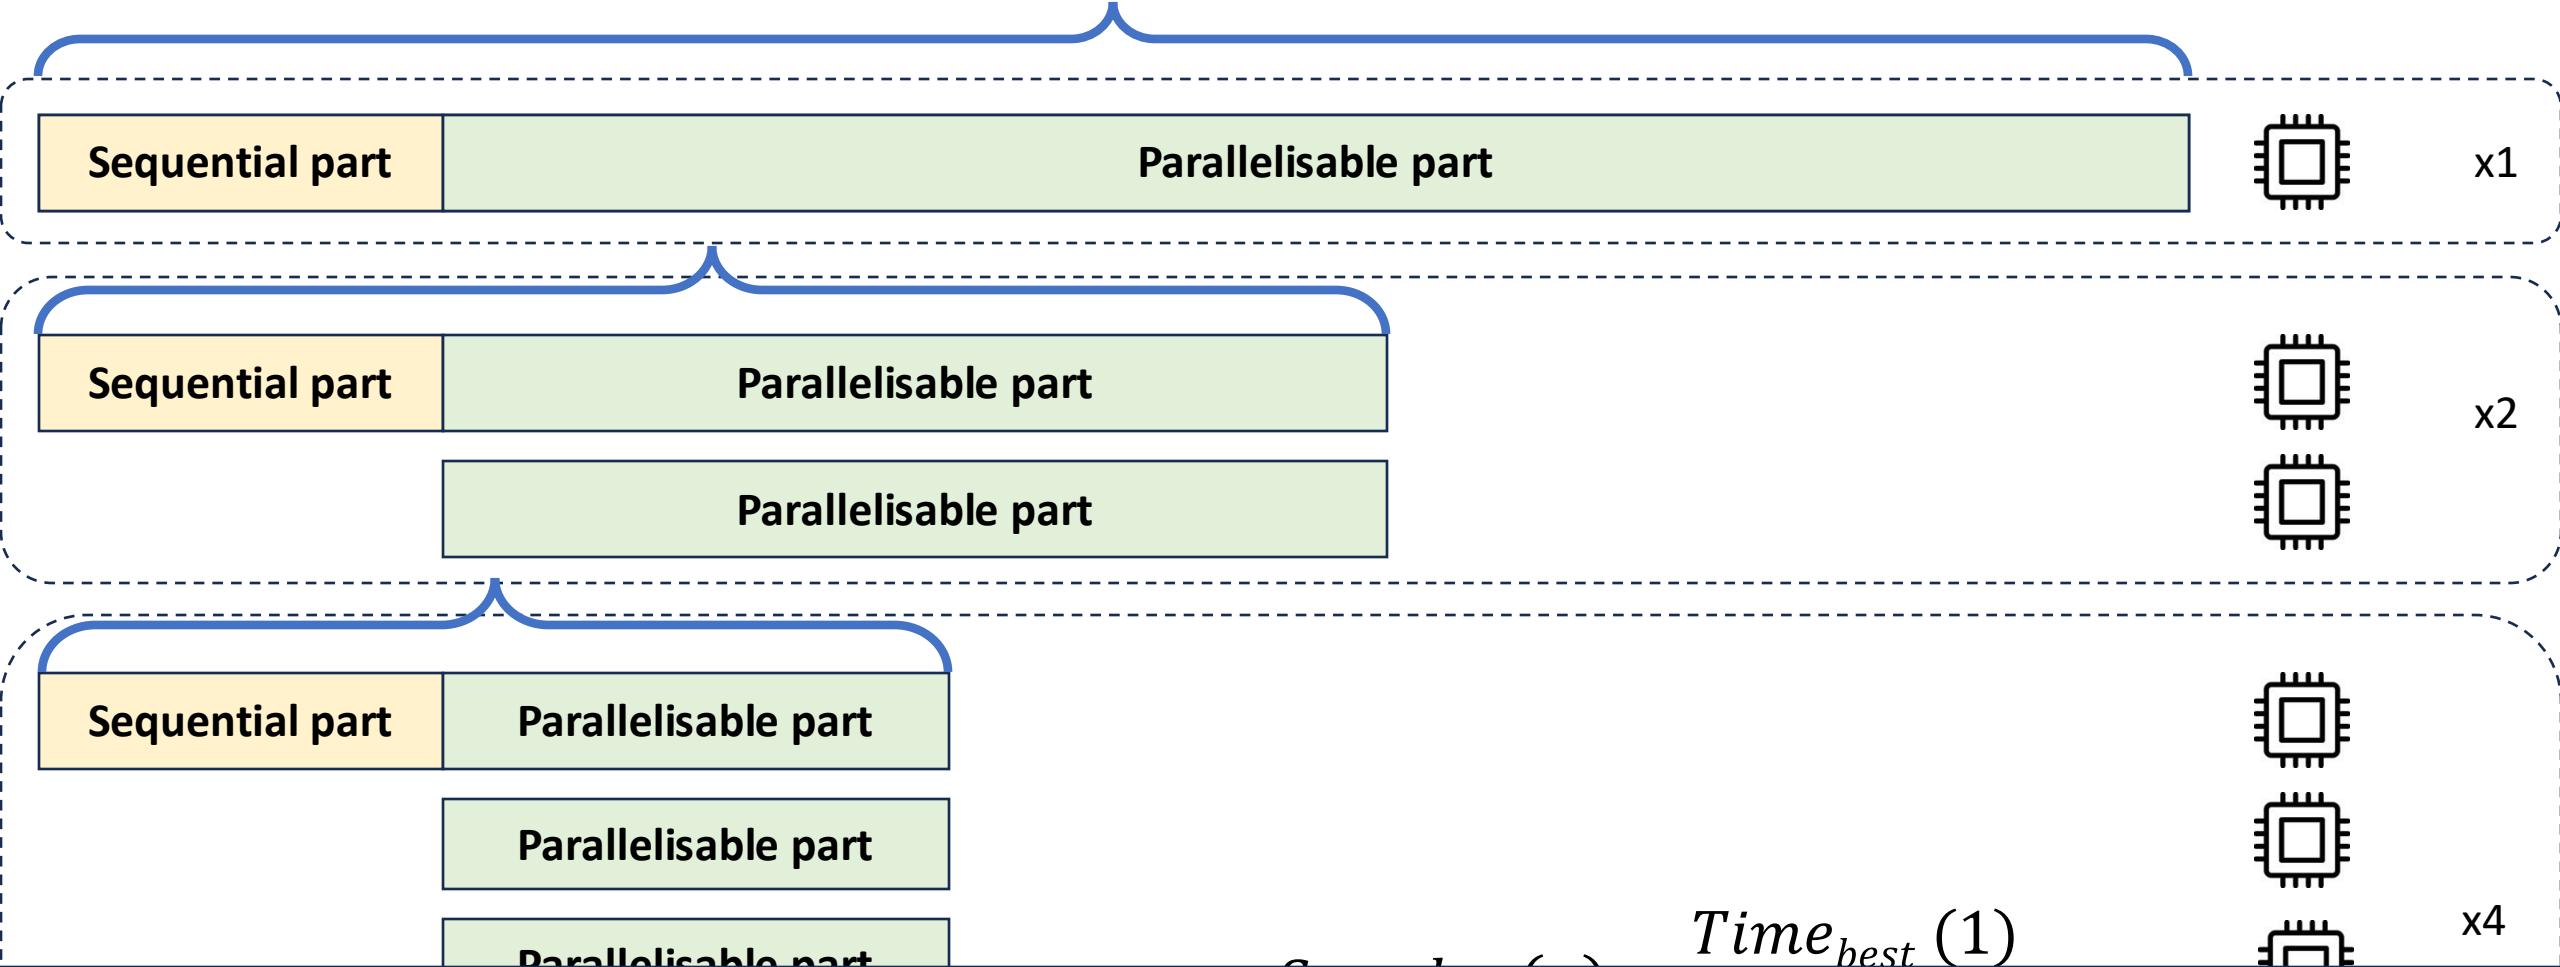
# How do threads speed up a sequential application?

- Can we speed up a data parallelisable sequential application using multithreads on a single core?
  - **No!** We need more cores to make concurrent threads in parallel.
- Can we speed up a task parallelisable sequential application using multithreads on a single core?
  - **Yes!** Avoid busy waiting in the sequential execution

# How do multi cores speed up a multithread application?

- For a **given problem size**, how will the program be **speed up** when **adding more processors(cores)**?
- For a **given multithread program**, how will it handle **bigger problem size** when **adding more processors(cores)**?

# How do multi cores speed a multithread application up?



Discussion: How will race conditions affect the speed-up?

# Amdahl's Law

- The speedup of an application when using multiple processors is limited by the app's sequential parts

- Thus:

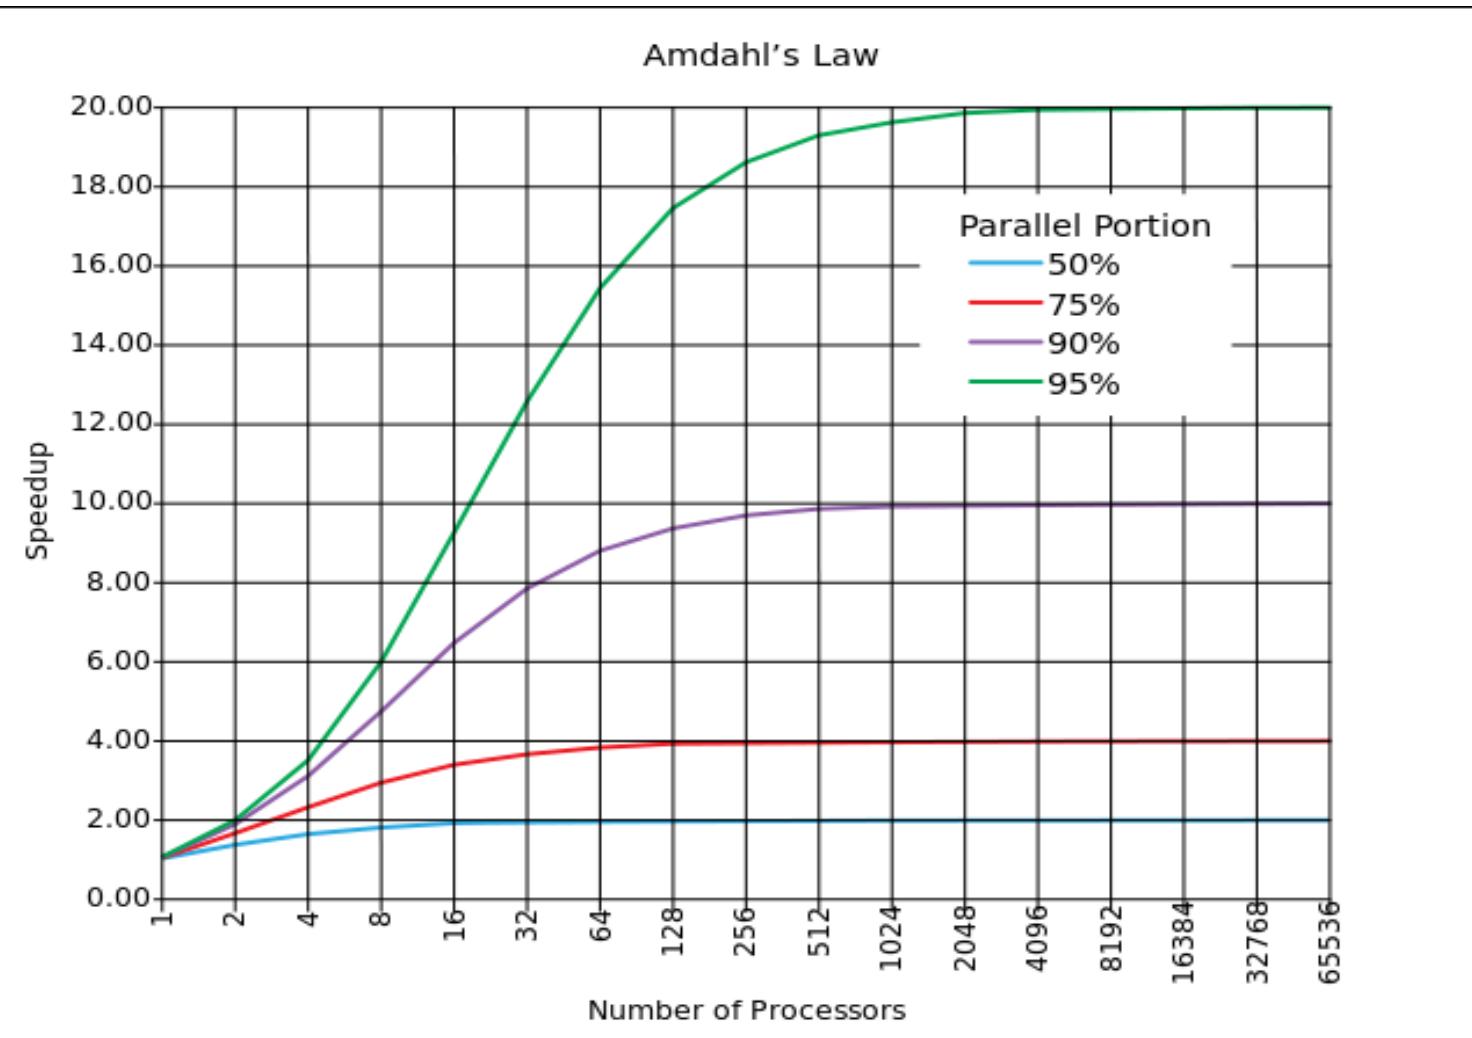
$$Time_{new} = Time_{sequential\_part} + \frac{Time_{parallelizable\_part}}{Number\ of\ processors}$$

- Or:

$$Speedup = \frac{1}{(1 - Fraction_{parallelizable\_part}) + \frac{Fraction_{parallelizable\_part}}{Number\ of\ processors}}$$

# Amdahl's Law in pictures

$$Speedup = \frac{1}{(1 - Fraction_{parallelizable}) + \frac{Fraction_{parallelizable}}{Number\ of\ processors}}$$



$$Speedup_{up\_limit} = \frac{1}{(Fraction_{sequential})}$$

# Efficiency

How well the application utilises the computational resources to achieve the speedup.

$$Efficiency(n) = \frac{Speedup(n)}{n} = \frac{Time_{best}(1)}{n * Time(n)}$$

# Efficiency

Efficiency can also be seen as the ratio between the “ideal performance” and the “actual performance”.

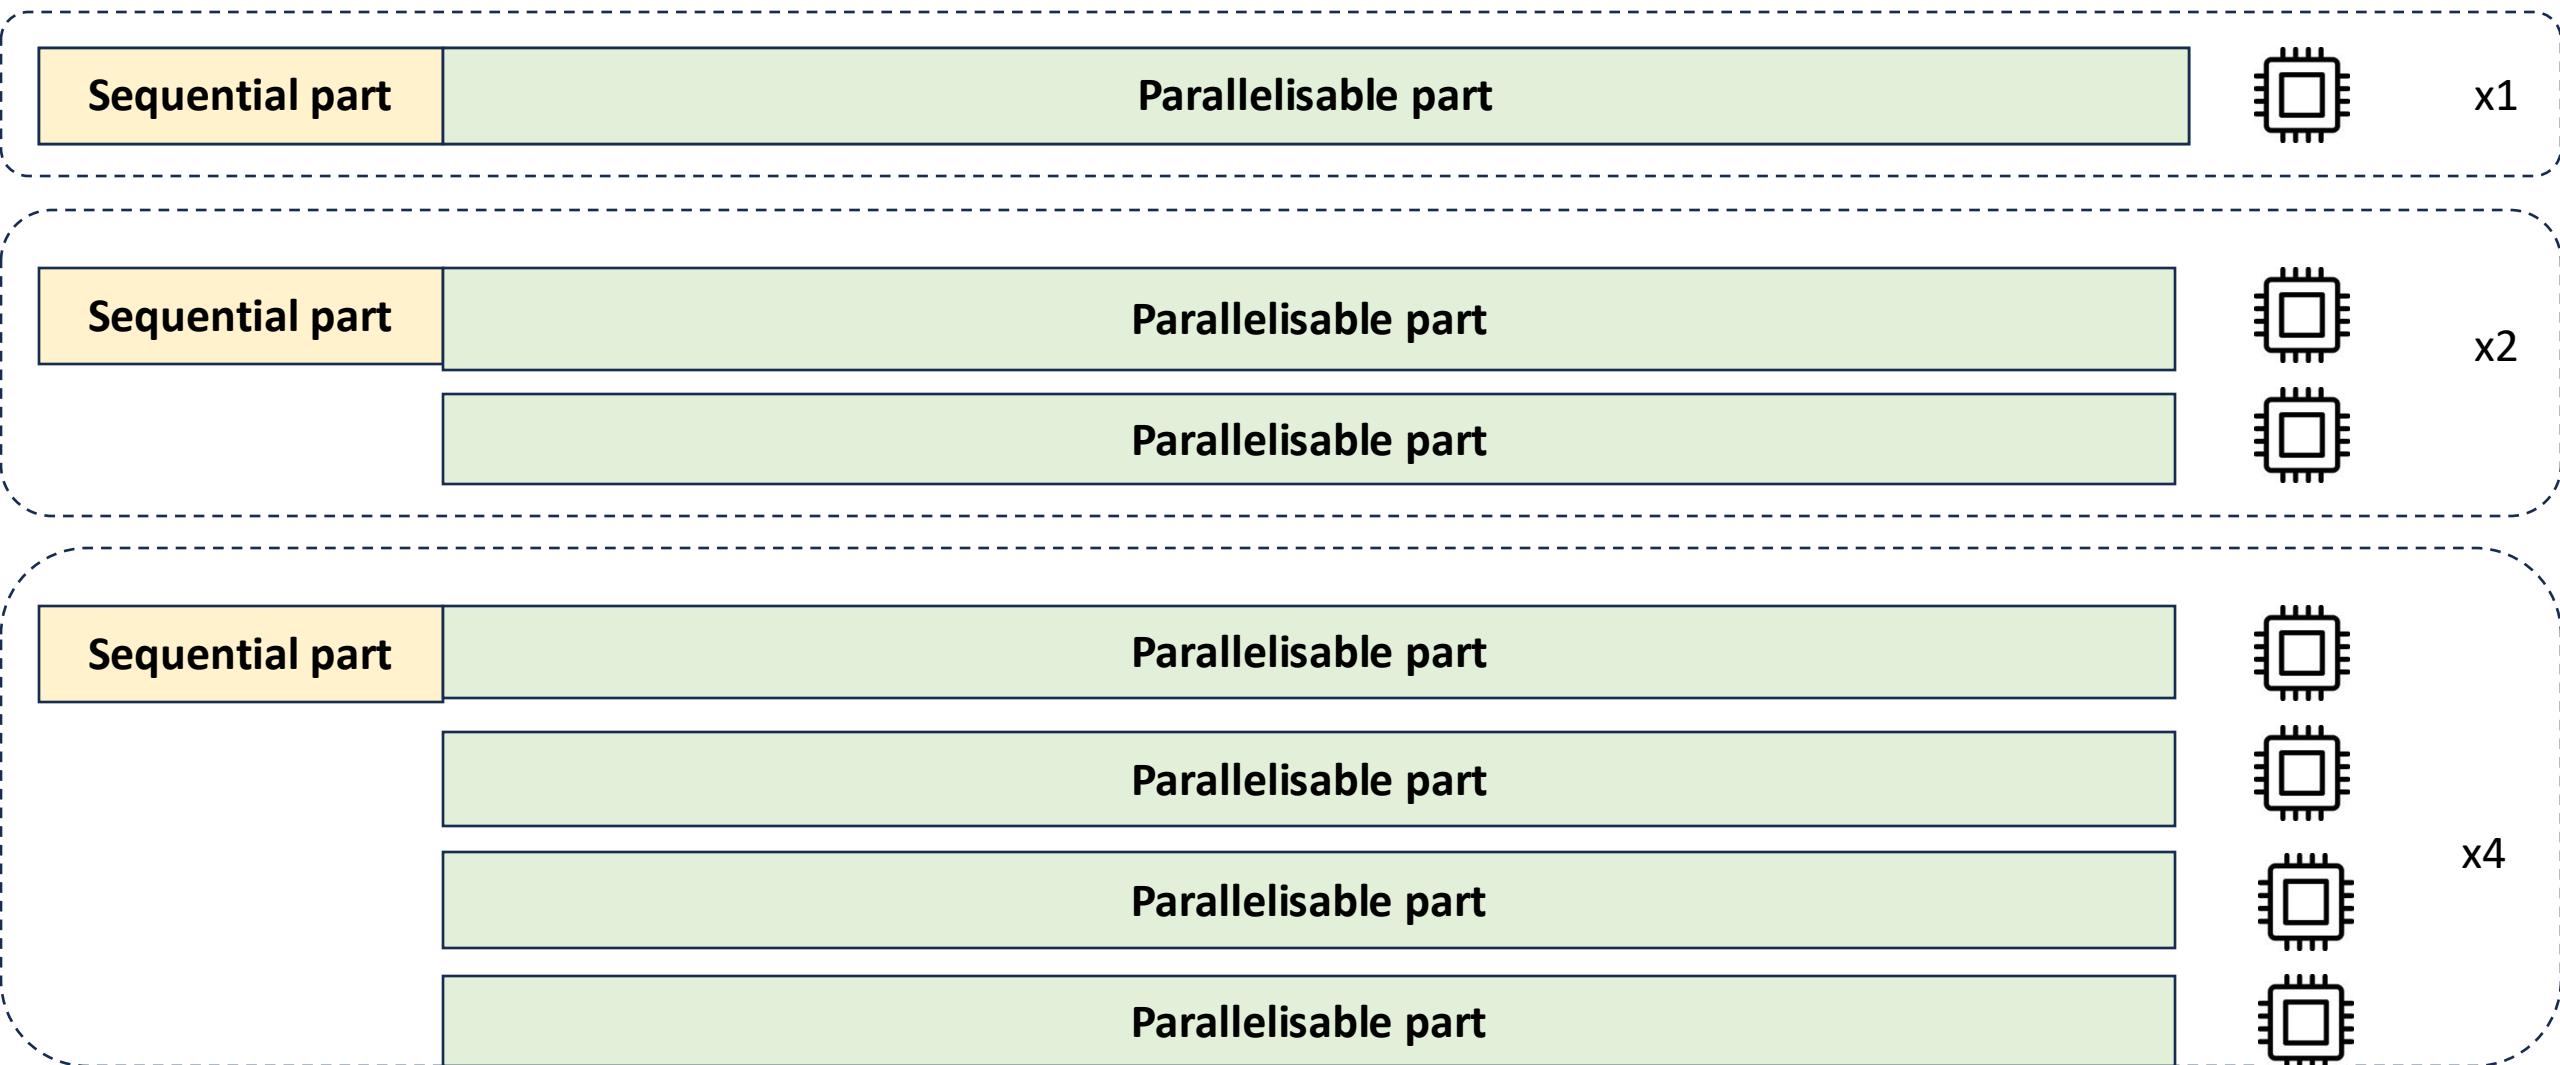
$$\begin{aligned} \text{Efficiency}(n) &= \frac{\text{Speedup}(n)}{n} = \frac{\text{Time}_{best}(1)}{n * \text{Time}(n)} \\ &= \frac{\text{Time}_{ideal}(n)}{\text{Time}(n)} \end{aligned}$$

$$\text{Time}_{ideal}(n) = \frac{\text{Time}_{best}(1)}{n}$$

# How do multi cores speed up a multithread application?

- For a **given problem size**, how will the program be **speed up** when **adding more processors(cores)**?
- For a **given multithread program**, how will it handle **bigger problem size** when **adding more processors(cores)**?

# How problem size changes with more cores?



# Gustafson's Law in pictures

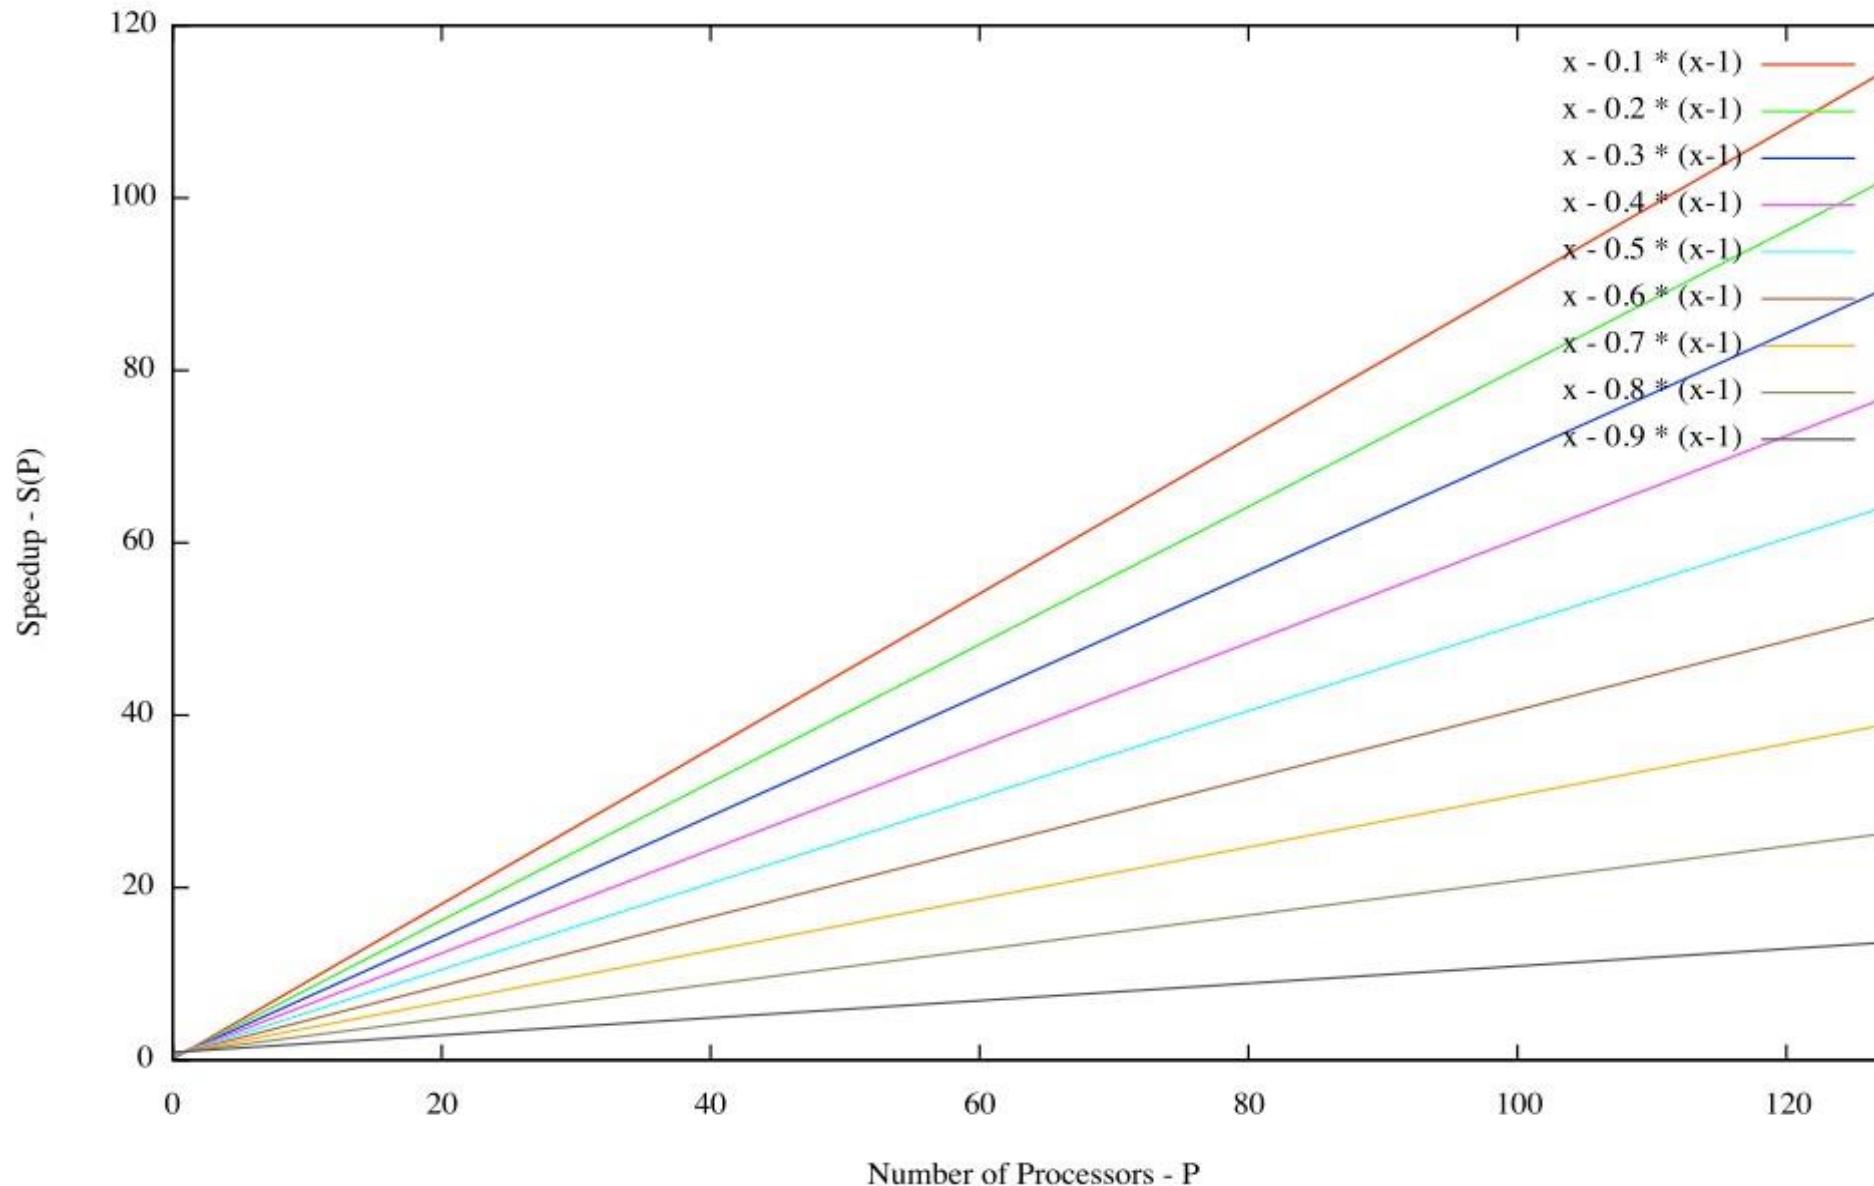
- Allows for scaling the problem size
- Gustafson's law considers the execution time of a parallel program using  $n$  processors with a sequential fraction  $s$  (with  $s$  in  $[0,1]$ )
- *Translates into:*

$$\text{Speedup} = \frac{\text{Fraction}_{\text{Sequential part}} + n * (1 - \text{Fraction}_{\text{Sequential part}})}{\text{Fraction}_{\text{Sequential part}} + (1 - \text{Fraction}_{\text{Sequential part}})}$$

$$\begin{aligned}\text{Speedup} &= \text{Fraction}_{\text{Sequential part}} + n * (1 - \text{Fraction}_{\text{Sequential part}}) \\ &= n - \text{Fraction}_{\text{Sequential part}} * (n - 1)\end{aligned}$$

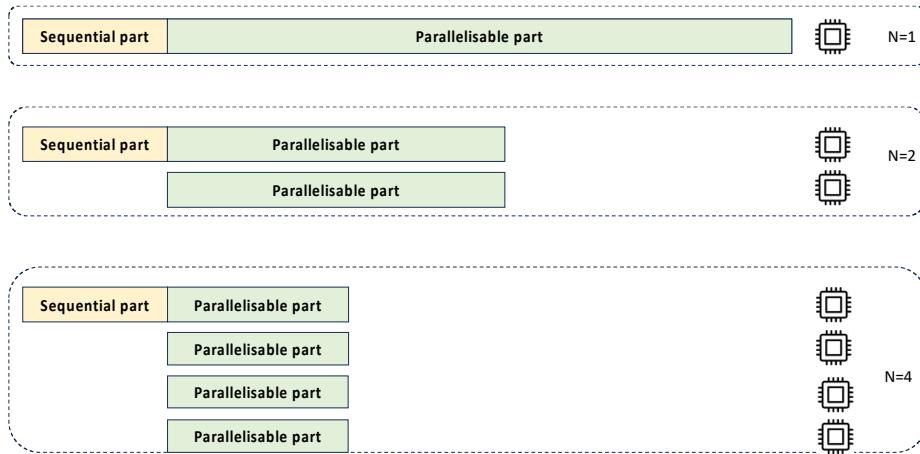
# Gustafson's Law in pictures

Gustafson's Law:  $S(P) = P - a * (P-1)$

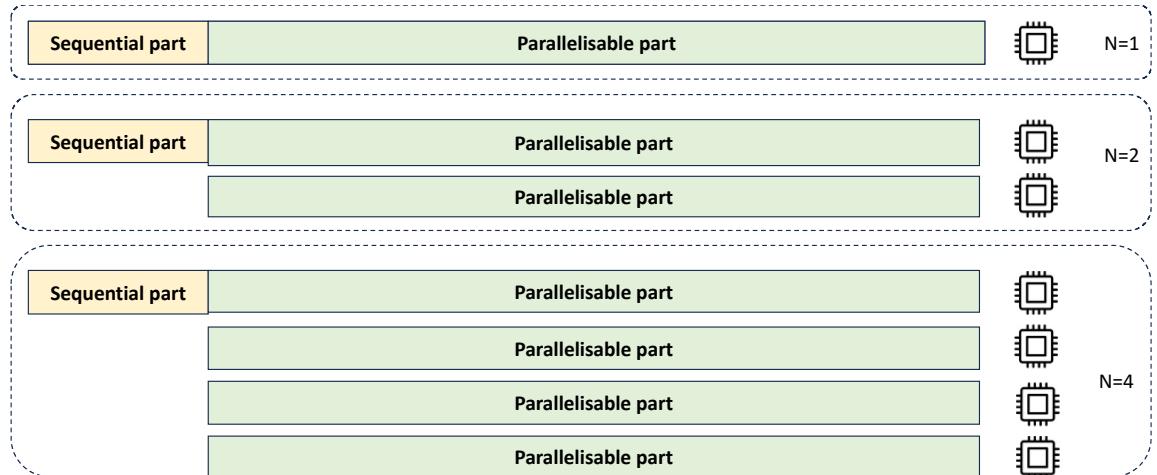


# What is the difference between Amdahl's and Gustafson's law?

- Amdahl's law



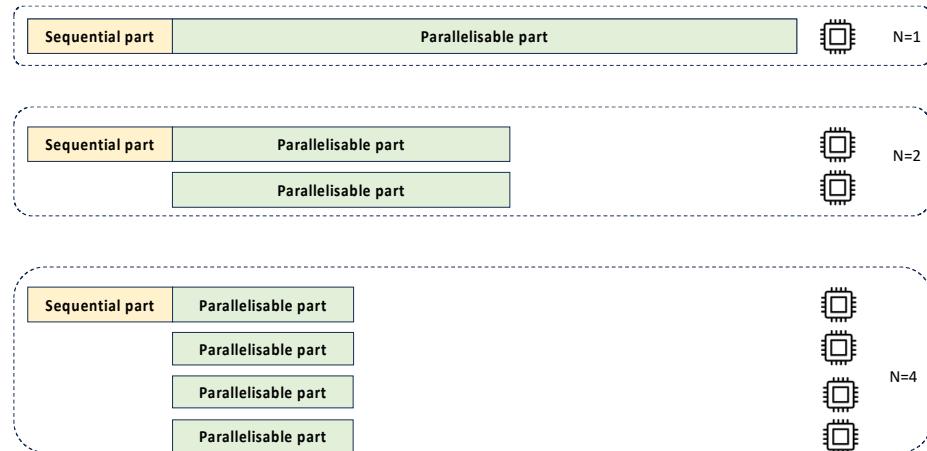
- Gustafson's law



# What is the difference between Amdahl's and Gustafson's law?

- Amdahl's law

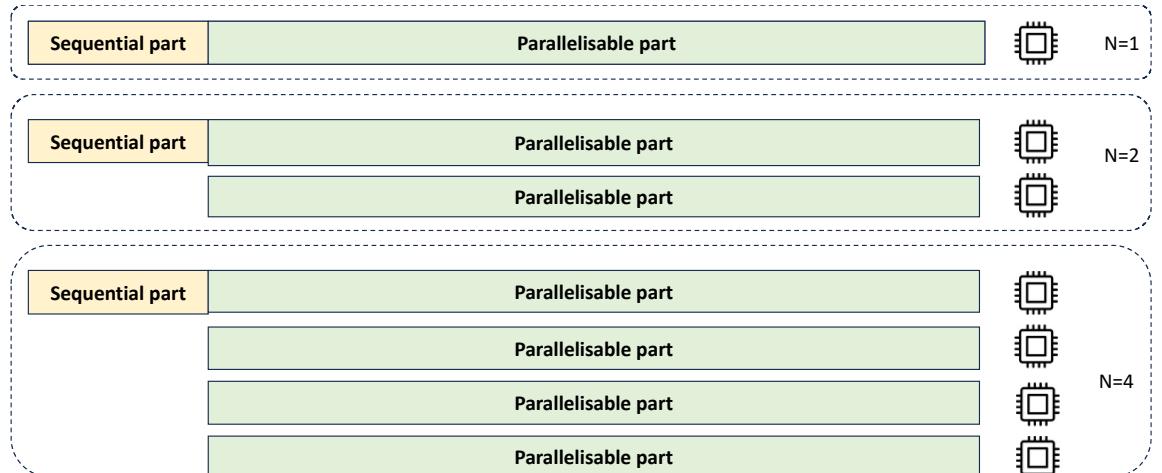
How does the *speedup* of a **fixed-size problem scale** when it is executed on **N cores**?



**Strong scaling**

- Gustafson's law

How does the *speedup* of a **changing size problem scale** when it is executed on **N cores**, but with **fixed size on each core**?



**Weak scaling**