

# Message passing Interface: derived data types in MPI

---

Adam Belloum

Source: Parallel Programming 2020: Lecture 11- MPI data types, virtual topologies, and performance pitfalls

# MPI data types: Why?

```
MPI_Bcast (&cfg.nx, 1, MPI_INT, ...);  
MPI_Bcast (&cfg.ny, 1, MPI_INT, ...);  
MPI_Bcast (&cfg.du, 1, MPI_INT, ...);  
MPI_Bcast (&cfg.nx, 1, MPI_DOUBLE, ...);  
MPI_Bcast (&cfg.it, 1, MPI_INT, ...);
```



Want to do:

```
MPI_Bcast (&cfg, 1, <type cfg>, ...);
```

```
MPI_Bcast (&cfg, 1, sizeof(cfg), MPI_BYTE, ...);
```

- Works in practice
- But not portable

# MPI data types: Why?

Example: Send column of a matrix (**non-contiguous** in C) ?

- Send each element alone ?
- Manually copy element into a contiguous buffer and sent it?
- New data type !!!
- Three step process

## I. **Construct** with

```
MPI_Type_* (MPI_Datatype *int);
```

## 2. **Commit** new data type with

```
MPI_Type_commit(MPI_Datatype *int);
```


## 3. **Deallocate**

```
MPI_Type_free(MPI_Datatype *int);
```

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29

# A flexible, vector-like type: **MPI\_Type\_vector**


```
MPI_Type_vector (int count,          2 // no. of blockint blocklength, 3 // no. of element/blockint stride,      5 // no. of elements b/w start of each blockMPI_data_Type oldtype, MPI_Datatype * new type);
```




```
MPI_Data_type nt;
MPI_Type_Vector(2, 3, 5, MPI_INT, &nt);
MPI_Type_commit(&nt);
< Do something >
MPI_Type_free(&nt);
```

# Caveat when using a type

- **Caution:**
  - **Concatenating** such a type in a **send** operation can lead to unexpected results
  - Count argument to send and others must be handled with care:



```
MPI_Send(buf, 2, nt, ...);
```



# Derived type size and extend


- Get the total size (in byte) of datatype in a message

```
MPI_Type_size(MPI_Datatype new_type, int *size);
```

- Get the
  - lower bound**
  - the **extent** Span from **1<sup>st</sup>** to the **last byte** in a Datatype


```
MPI_Type_get_extent(  
    MPI_Datatype newtype,  
    MPI_Aint *lb,  
    MPI_Aint *extent);
```

- In MPI to change the extent of a datatype
  - Using **lb\_marker** and **ub\_marker**
  - Does not affect the size of the datatype
  - Does affect the outcome of the replication of this data type



# Sending a column of a Matrix in C

Raw-major datatype in C → cannot use plain array



```
double matrix[30];
MPI_Datatype nt;


// count = nrows, blocklength = 1,
// stride = ncols
MPI_Type_vector(nrows, 1, ncols,
                 MPI_DOUBLE, &nt);
MPI_Type_commit(&nt);

// send column
MPI_Send(&matrix[1], 1, nt, ...);

MPI_Type_free(&nt);
```

# A sub-array type: **MPI\_Type\_Create\_subarray**

```
MPI_Type_create_subarray (
    int dims,           dim of the initial array
    int ar_sizes[],   array - sizes of array (in each dim)
    int arsizes[],    array - sizes of subarray (in each dim)
    int ar_starts[],  Start indices of subarray
    int order,
    MPI_datatype oldtype
    MPI_Datatype * new type);
```




0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29

Diagram illustrating a 6x5 grid of numbers from 0 to 29. A brace on the left is labeled 'nrows' and a brace at the bottom is labeled 'ncols'.

- Order:
  - Row-major: **MPI\_ORDER\_C**
  - Column-major: **MPI\_ORDER\_Fortran**

# Example for a sub-array type: bulk of a matrix

```
Dims          2
ar_sizes     {nclos, nrows}
ar_subsizes  {nclos-2 , nrows-2}
ar_starts    {1, 1}
Order         MPI_ORDER_C
Oldtype       MPI_INT
```




```
MPI_create_Subarray (dims, ar_sizes, ar_subsizes, ar_starts,
order, oldtype,&nt);

MPI_Type_commit (&nt);
//Use nt ...
MPI_Send (&buf[0], 1, nt, ...); //etc
MPI_type_free(&nt)
```


# Most flexible type: `MPI_Type_create_struct`

```
MPI_Type_create_struct (
    int count,           // number of type composing the struct
    int bloc_length[],  // the length of type composing the struct
    int Aint displs[],  //
    MPI_datatype types[],           Count. =2
    MPI_Datatype * newtype);
```



The contents of `displs` are either

- the displacements in `bytes` of the block bases
- or `MPI addresses`



# How to obtain and handle addresses?

```
MPI_Get_address(const void *location, MPI_Aint *address);  
MPI_Aint MPI_Aint_diff(MPI_Aint addr1, MPI_Aint addr2)  
MPI_Aint MPI_Aint_add(MPI_Aint base, MPI_Aint disp)
```

## Example

```
Double a[1000];  
MPI_Aint a1, a2, disp;  
MPI_Get_address(&a[0], &a1);  
MPI_Get_address(&a[50], &a2);  
Disp = MPI_Aint_diff (a2,a1)
```

# Derived data type: summary

- A flexible tool for communicate complex data structures in MPI
- Most important calls
  - `MPI_Type_vector` (second simplest)
  - `MPI_Type_create_subarray`
  - `MPI_Type_create_struct` (most advanced)
  - `MPI_Type_commit/MPI_Type_free`
  - `MPI_Get_address,`
  - `MPI_Aint_add, MPI_Aint_diff`
  - `MPI_Type_get_extent, MPI_Type_size`
- Other useful features
  - `MPI_Type_contiguous`, `MPI_Type_indexed`, ...
- **Marching rule:**
  - **send** and **receive** match if specified **basic datatypes** match **one by one**,
  - Correct displacement at receiver side are automatically matched to the corresponding data items


# Message passing Interface: Virtual (Cartesian) topologies in MPI

---

# A convenient process naming scheme for multi-dimension problems

- Convenient process naming
  - Naming scheme to fit the communications pattern
  - Simplifies writing of code
  - Can allow MPI to optimize communications
- Let MPI map ranks to coordinates
- User: map array segments to ranks


Example: distribute 2-D array of  
400x 3000 elements  
equally on 12 ranks  
4000 elements



# Creating a Cartesian communicator

- Create a **new communicator** attached to cartesian topology

```
MPI_Cart_create(  
  MPI_Comm oldcomm,  
  int ndims,                      number of dimensions  
  int dims[],                    array with ndims elements, dims[i] specifies  
  int periods[],                  the number of ranks in dimension i  
  int reorder,                   array with ndims elements, periods[i]  
  MPI_Comm *cart_somm           specifies if dimension i is periodic  
);  
                                allow rank of oldcomm to have a different  
                                rank in cart_comm  
  
                                ndims = 2  
                                dims[0]=4  
                                dims[1]=3  
                                Periods=0/1
```




# Cartesian topology example

**ndims** = 2

**dims.** = { 4, 3 }

**Periods** = { 1, 0 }

**reorder** = 0



# Cartesian topology service functions

- Retrieve rank in new Cartesian communicator ("who am I in the grid ?")

```
MPI_Comm_rank(cart_comm, int * cart_rank);
```

- Map rank → coordinates ("where am I in the grid ?")

```
MPI_Comm_coord(cart_comm, rank, int maxdims, int coords[]);
```

**rank**: any rank is part of the cartesian comm `cart_comm`

**coords**: array of `maxdims` elts, receives the coordinates for `rank`

- Map coordinates → rank (Who is at the position ?")


```
MPI_Cart_rank(cart_comm, int coords[], int *rank);
```

**coords**: coordinates; if periodic in direction *i*,

`coords[i]` are automatically mapped int the valid range, else they are erroneous

# Example

- Example 12 processes arranged in 4X3 grid
- Column-major numbering
- Process coordinates begin with 0



# Next-neighbour communication

- Sending/receive from **neighbours** is a typical task in cartesian topologies

```
MPI_Cart_shift (
    cart_comm,
    direction,    direction: dimensions to shift
    disp,         offset to shift:
                  >0 shift in positive direction,
                  >0 shift in negative direction
    int *src_rank,
    int *dest_rank );
```

retuned ranks as input (argument) of `MPI_Sendrec*` calls

# Next-neighbour communication

Example : 4X3 progress grid, **period on 1st dimension** (dim0 - horizontal),  
each process has an int value which get shifted


```
MPI_Cart_shift (cart_comm, 0, 1, &src, &dest);
```

```
MPI_Sendrecv_replace(&value, 1, MPI_INT, dst, 0, src, cart_comm, ...);
```



```
MPI_Cart_shift (cart_comm, 1, 1, &src, &dest);
```

```
MPI_Sendrecv_replace(&value, 1, MPI_INT, dst, 0, src, cart_comm, ...);
```



For non-periodic topologies

- `MPI_PROC_NULL` is returned in boundaries

# Message passing Interface: typical performance pitfall in MPI


---

# Implicit serialization and synchronization

Common performance pitfall with MPI


Consider linear shift in an **open chain**, e.g.,

- **each rank** in the chain issues:
  - `MPI_Send(..., rank+1, ...)`;
  - `MPI_Recv(..., rank-1, ...)`;
- **First and last rank** call → `MPI_Send` and `MPI_Recv` only
- There is no danger of deadlocks
- **But** performance depends on implementation-specific parameters in the MPI library: `MPI_Send`




# Implicit serialization and synchronization

- Best case scenario:
  - `MPI_Send` operates in a buffered send mode
- `MPI_Send` returns after message is copied to a system buffer
- Send/Receive operations can be overlapped on nonblocking, bidirectional networks



# Implicit serialization and synchronization


- **Worst case scenario 1:** Synchronous send using
  - the **rendezvous protocol**



Rendezvous: Send blocks until complete message has been transferred!

# Implicit serialization and synchronization


- **Worst case scenario 2:** Synchronous send using the **eager protocol**
- **Eager:**
  - Message may be transmitted to receiver without a matching receive issued.
  - Data is put in a local system buffer at receiver side



# Implicit serialization and synchronization

## Better implementation alternatives:

- Reverse order of send/receive calls on even and odd ranks
- Use non-blocking `MPI_Isend`/`MPI_Irecv` pairs. Multiple outstanding/open communication requests allow flexible scheduling. Also: potential for asynchronous data transfer.
- Use `MPI_Sendrecv` or `MPI_Sendrecv_replace`: Simple coding with flexible message scheduling; but no asynchronous transfer



# Observing Synchronization Delays

- 3 processors sending data, with one sending a short message and another sending a long message to the same process:



# Network contention

Contention on network level may occur:

- Multiple processes on a node use the **network interface** at the same time.
  - (Network bandwidth per process decreases linearly if a single process can already achieve full network bandwidth)
- **Network topology** is not fully non-blocking,
  - i.e. bisection bandwidth/compute node decreases with increasing compute nodes.
- Non-optimal routing:
  - Even for full non-blocking networks contention may occur on **internal network links**
- **MPI\_Alltoall** – Communication pattern most vulnerable:
  - Every process wants to talk to everyone else at the same time (imagine 300.000 processes do that)

# Non-blocking but non-asynchronous MPI calls

- **Intention:**
  - Use “non-blocking” calls for communication overlap:  
`MPI_ISEND(A, ..., &request);`  
    ...do **useful work** and do not modify **A**  
`MPI_WAIT(&request, ...);`
  -
- **Perfect world:**
  - “**useful work**” takes longer than communication
  - Nonblocking MPI call **implements asynchronous data transfer**
- **However,**
  - the MPI standard **does not guarantee** asynchronous transfer

# A simple test for asynchronicity

- Do some work for some configurable time (`delay`) while a message of `count` bytes is received or sent

```
if(rank==0) {  
    t = MPI_Wtime();  
    MPI_Irecv(buf, count, MPI_BYTE, 1, 0, MPI_COMM_WORLD, &req);  
    do_work(delay);  
    MPI_Wait(&req, &status);  
    t = MPI_Wtime() - t;  
} else {  
    MPI_Send(buf, count, MPI_BYTE, 0, 0, MPI_COMM_WORLD);  
}  
printf("Overall: %lfDelay: %lf\n", t, delay);
```


- If overlap occurs, overall time will be constant w.r.t. the delay up to a certain point

# Results

- Intel MPI 2019 update 8
- Message size = 240 MB
- Default behaviour: no overlap
- `I_MPI_ASYNC_PROGRESS=1`:  
// full overlap observed!

General remarks: a **moving target**;

- MPI implementations change all the time
- **Depends on** inter-/intranode, message size, network layer,...
- MPI implementations provide tuning knobs



# MPI performance pitfalls summary

- Always observe **possible synchronizing properties** of MPI calls
- Assume the worst possible behaviour
- Even if a deadlock does not occur, performance may be severely impacted
- **Network contention** is a fact
- Most network connections can be saturated with a single connection
- Most large-scale networks are not entirely non-blocking for cost reasons
- **Non-blocking** MPI communication is **not necessarily asynchronous**
- Study possible tuning knobs